

Hyperparameter Tuning: Search

Hyperparameter tuning is the process of finding the optimal combination of hyperparameters for a machine learning model. A hyperparameter search consists of four key components:

1. **Hyperparameter space:** This defines the range of possible values for each hyperparameter you want to tune. For example, for a random forest, this might include the number of trees, maximum depth, minimum samples per leaf, etc.
2. **Sampling method:** This determines how candidate hyperparameter combinations are selected from the space. Common approaches include:
 - o Grid search (exhaustively tries all combinations)
 - o Random search (randomly samples combinations)
 - o Bayesian optimization (intelligently selects promising combinations based on previous results)
3. **Cross-validation scheme:** This is how you evaluate each candidate hyperparameter set. Typically k-fold cross-validation is used, where the data is split into k parts, with each part serving as the test set once while the others are used for training.
4. **Performance metric:** This is what you're trying to optimize (either minimize or maximize). The choice depends on your problem - accuracy, precision, recall, RMSE, etc.

Metrics

Metrics are quantitative measures used to evaluate model performance. They vary by problem type:

Classification Metrics (for categorical outcomes):

- **Accuracy:** Percentage of correct predictions
- **Precision:** Percentage of positive predictions that were correct
- **Recall:** Percentage of actual positives that were correctly identified
- **F1-score:** Harmonic mean of precision and recall

- **ROC-AUC:** Area under the receiver operating characteristic curve

Regression Metrics (for continuous outcomes):

- **Mean Squared Error (MSE):** Average of squared errors
- **Root Mean Squared Error (RMSE):** Square root of MSE
- **Mean Absolute Error (MAE):** Average of absolute errors
- **R-squared:** Proportion of variance explained by the model

In Scikit-learn:

Scikit-learn provides these metrics in its `metrics` module. They're used by:

1. Importing the relevant metric function
2. Passing true values and predictions to the function
3. For hyperparameter tuning, the metric is specified in the search object (like `GridSearchCV`)

Creating Custom Metrics:

You can define your own metrics when standard ones don't capture your specific needs.

This involves:

1. Creating a function that takes true values and predictions
2. Implementing your custom calculation
3. Ensuring it returns a single numerical value (higher=better or lower=better)
4. Using it in your model evaluation or hyperparameter search